

The Limits of Diagnostic Bug Taxonomies for Failed Human and LLM Competitive-Programming Submissions

Anonymous Author(s)

Submitted for review

Abstract—Bug taxonomies are attractive for explaining failed LLM code, but before comparing human and model failures they raise a measurement question: can labels assigned from the usual evidence—problem statement, buggy code, and verdict—support distributional claims? We study this question on 14,182 sandbox-confirmed buggy C++ submissions for 107 post-cutoff CodeContests problems: human submissions, a cutoff-aligned proxy model, a modern contamination-uncontrolled model, and a self-reflection repair loop. We reuse the 32-leaf taxonomy of Wei et al., collect a 375-bug doubly annotated gold set, validate an LLM judge against it, and then audit a separate 120-item subset with stronger evidence: first failing test, actual/expected output, and a reference solution. The measurement result is central. The LLM judge reaches substantial agreement with the gold ($\kappa \approx 0.76$), but over-applies the broad design family by 6–9 points, the same scale as many full-corpus human–AI shifts. More importantly, stronger evidence changes 55.0% of family labels. Thus common-evidence taxonomy labels are useful diagnostic proxies, not demonstrated root causes. Under that bound, only a coarse profile is defensible: human and AI failures all concentrate in mathematics and design, while low-level coding failures remain minority modes. Label-free dynamics are clearer: the cutoff-aligned proxy remains near the floor, the modern contaminated reference is stronger but difficulty-sensitive, and public-example self-reflection repairs only 4.0% of proxy-model bugs—matching a same-budget resampling baseline and fixing no runtime or time-limit failures. We release the corpus, prompts, cached outputs, labels, and analysis scripts.

Index Terms—large language models, competitive programming, bug taxonomy, empirical study, LLM-as-a-judge, data contamination

I. INTRODUCTION

Code-generating large language models (LLMs) have moved from toy benchmarks to competitive programming, a domain demanding precise algorithmic reasoning, careful edge-case handling, and strict conformance to input/output (I/O) formats [1], [2]. Most work measures *whether* models solve such problems, via $\text{pass}@k$ on benchmarks like HumanEval, MBPP, and APPS [3]–[6]. Far less is known about *how* models fail, and almost nothing about whether their failures resemble those of human programmers solving the same problems.

The question matters for tooling and science. If AI and human bugs cluster in the same places, the existing ecosystem of human-oriented review, linters, and education transfers to AI-assisted development; if they differ systematically, that ecosystem must be re-pointed. Measuring this is hard for three reasons. First, benchmarks suffer from *data contamination*:

problems and editorial solutions appear in pre-training data, so apparent failures may reflect memorization rather than reasoning [2], [7]. Second, a fair human–AI comparison needs the *same* problems and a *common* bug vocabulary. Third, labeling tens of thousands of buggy programs into a fine-grained taxonomy is beyond manual reach, pushing studies toward automated (LLM) labeling whose reliability is itself unestablished.

We address all three. We adopt the recent bug taxonomy of Wei et al. [2]—six general-error and six algorithm-specific families, 32 leaves (Table I)—and bound contamination risk by using the test split of CodeContests [1]: 107 Codeforces problems first published after the September-2021 cutoff of `gpt-3.5-turbo` [8]. On these fixed problems we compare four bug sources: human submissions; a cutoff-aligned proxy model; a modern model whose training window covers the problems (a present-day performance reference, not a clean source); and a self-reflection loop on the clean model. Every submission is judged in a hardened sandbox; two independent annotators label a stratified 375-bug gold set; and an LLM judge labels the full corpus, which we validate against the human gold and the human inter-annotator ceiling.

We answer four research questions, ordered by the strength of evidence needed for the claim:

- **RQ1:** How stable are taxonomy labels under different labelers and stronger failure evidence?
- **RQ2:** What bug-profile claims remain defensible after those measurement limits?
- **RQ3:** How do capability, contamination risk, and problem difficulty shape label-free failure dynamics?
- **RQ4:** What does public-example self-reflection repair, and which failures persist?

Our findings reframe the comparison. (1) A validated LLM judge is adequate for coarse triage but not for fine human–AI separation: it reaches substantial agreement with the gold while inflating design labels by 6–9 points. (2) The root-cause audit is a stronger warning: first-failing-test/reference evidence changes 55% of family labels, so common-evidence labels must be read as diagnostic proxies. (3) Within that bound, the stable bug-profile result is coarse: mathematics and design dominate every source, and low-level coding errors are minority modes. Full-corpus family shifts are visible but small and exploratory. (4) Label-free dynamics are more interpretable:

the proxy model remains near the floor, while the modern contaminated reference is stronger but difficulty-sensitive. (5) Self-reflection with only public-example feedback repairs 4% of proxy-model bugs, matching a same-budget resampling baseline, with a sharp difficulty collapse (7.1%→1.4%) and no runtime or time-limit fixes.

Contributions: (i) a paired, contamination-aware artifact for failed-submission measurement (14,182 bugs, 107 identical problems, 375 doubly-annotated gold items); (ii) a measurement protocol—human gold, power analysis, problem-level bootstrap, LLM-judge bias analysis, and a first-failing-test root-cause audit—that separates defensible coarse claims from fine claims the instrument cannot support; (iii) empirical evidence that common-evidence taxonomy labels are often not root-cause labels; and (iv) label-free and taxonomy-aware self-reflection analyses. All artifacts are released.

II. BACKGROUND AND RELATED WORK

A. LLMs for code and competitive programming

Transformer language models [9], [10] pre-trained on source code now span open and proprietary families [3], [8], [11]–[16], and a broad literature surveys their software-engineering uses [17]. Competitive programming is repeatedly singled out as the hardest target because it couples non-trivial algorithmic reasoning with strict resource and I/O-format constraints [1], [2]: AlphaCode [1] reached competition-level generation only through massive sampling and filtering. Most evaluation reports *whether* a model succeeds—pass@ k on HumanEval, MBPP, APPS, and successors [3]–[5], [18], or repository-level tasks like SWE-bench [19], drawing on code corpora and contest archives [20]–[22]—and rigorous re-testing shows such pass rates are often optimistic [6]. We instead study *how* models fail, and how that compares to humans on identical problems, a question these correctness-oriented benchmarks do not address.

B. Data contamination

Because contest problems and editorial solutions are public, they routinely enter pre-training corpora and inflate benchmark scores; a growing line of work measures and mitigates this [23]–[26]. LiveCodeBench [7] adopts rolling, recent problems and reports strong models near zero on the hardest unseen tasks, and Wei et al. [2] run an explicit contamination audit. We take the same stance structurally: the CodeContests test split post-dates the gpt-3.5-turbo cutoff (a clean model), and we add a modern model as a deliberately contamination-uncontrolled contrast rather than a clean measurement.

C. Bug taxonomies and empirical bug studies

Defect classification has a long lineage—Orthogonal Defect Classification [27], empirical bug characterizations [28]–[30], and curated fault benchmarks [31], [32] used to study repair [33]—almost all on *human* code; security studies of AI code (e.g., Copilot [34]) examine vulnerabilities rather than a general bug taxonomy. Closest to us, Wei et al. [2] recently proposed the taxonomy we adopt and hand-labeled

75 incorrect programs from a single model with strong inter-annotator agreement ($\kappa = 0.86$), using *no* automated labeler and providing *no* human baseline. We reuse their taxonomy verbatim but differ in kind as much as scale: we label **14,182** bugs (a $\sim 190\times$ larger corpus) from *four* sources—including *human* submissions to the *same* problems—under explicit cutoff alignment and contamination caveats, with an LLM judge whose reliability we validate against a 375-bug doubly-annotated human gold. The human baseline is what makes a human–AI comparison possible at all, and the scale is what forces the measurement questions—statistical power, within-problem clustering, and labeler granularity—that a 75-program, single-model, AI-only study never has to confront.

D. Automated repair and self-correction

Automated program repair has matured from search- and learning-based methods [32], [33] surveyed broadly [35], [36] to LLM-based repair [37]. For models repairing their *own* output, Self-Refine [38] critiques and revises with self-feedback, self-debugging [39] and Reflexion [40] add execution or verbal feedback, and Olausson et al. [41] question whether self-repair is a “silver bullet.” We instantiate a reflection loop with the public example-test results a contestant actually sees, and ask whether repair changes the *type* of remaining bugs rather than only the pass rate.

E. LLM-as-a-judge and agreement methodology

Using an LLM to score or label outputs is now widespread [42], [43], but it carries known biases [44] and its reliability for fine-grained, domain-specific taxonomies is rarely audited. We treat the judge as a measurement instrument: we size the human gold by Cochran’s proportion formula [45], [46], report chance-corrected agreement and its interpretation [47]–[51], and follow empirical-software-engineering guidance [52], [53].

F. The taxonomy we adopt

We do not propose a taxonomy; we reuse Wei et al.’s [2] verbatim (Table I) so our labels are comparable to theirs. It has two top levels. *General Errors* (GE) describe *how* a program is wrong regardless of the algorithm: design/wrong-algorithm (GE1), boundary and off-by-one (GE2), condition and operator mistakes (GE3), data-type issues such as integer overflow (GE4), syntax/compile errors (GE5), and input/output formatting (GE6). *Algorithm-specific Errors* (AE) describe *which* algorithmic technique was misapplied: mathematics (AE1), greedy (AE2), graph (AE3), recursion/divide-and-conquer (AE4), dynamic programming (AE5), and search (AE6). Each of the 12 families is split into leaves—e.g. GE1 separates *Incorrect Algorithm*, *Misunderstanding Requirements*, and *Inefficient design*—for 32 leaves total. Labeling is multi-label, since one program can exhibit several independent errors. The GE/AE split is the axis our analysis returns to: GE5/GE4 are coding errors a tool can catch, whereas GE1 and the AE families are reasoning errors that require understanding the problem.

TABLE I
THE TAXONOMY OF WEI ET AL. [2]: 6 GENERAL-ERROR (GE) AND 6 ALGORITHM-SPECIFIC (AE) FAMILIES, 32 LEAVES.

Code	Family (#leaves)
GE1	Design: wrong algorithm, misread requirement (4)
GE2	Boundary: edge cases, off-by-one/indexing (2)
GE3	Condition: predicates, operator/precedence (2)
GE4	Data type: overflow, implicit conversion (2)
GE5	Syntax: compile errors, language misuse (2)
GE6	Input/Output: parsing, output formatting (2)
AE1–AE6	Math, greedy, graph, recursion/D&C, DP, search (18)

III. STUDY DESIGN

Figure 1 summarizes the pipeline: build a cutoff-filtered benchmark; collect human and AI buggy submissions; confirm verdicts in a sandbox; label a human gold sample and scale with an audited LLM judge; analyze.

A. Benchmark, sandbox, and submissions

We use the CodeContests [1] *test* split: Codeforces problems first published after 2021-09-21, conservatively after gpt-3.5-turbo’s September-2021 cutoff [8]. We treat this cutoff as necessary but not sufficient—the exact training window of the -0125 snapshot is not publicly documented per-release—so we add a behavioral sanity check: the clean model’s acceptance is low and *flat* across difficulty ($\leq 8\%$; Sec. IV-C), which is consistent with solving from the statement rather than recalling seen tasks. We also run a lightweight code-similarity audit against the stored human accepted solutions for each problem: none of the proxy model’s 64 accepted outputs is whitespace-normalized identical to a stored human accepted solution, and the maximum token 5-gram Jaccard similarity is 0.47 (median 0.20). This does not prove absence of training exposure, but it rules out verbatim or near-verbatim replay of the accepted solutions available in our corpus. Throughout, “clean” means this cutoff-aligned proxy condition, not a provider-certified contamination-free model. Every cutoff-aligned taxonomy claim rests on this model; we make none on the modern model. We keep a problem only if it is judgeable (≥ 1 correct human solution confirmed AC and ≥ 10 incorrect ones confirmed non-AC), leaving **107 problems** (ratings 800–3500). The set is deliberately hard-skewed—25 problems rated ≤ 1199 , 12 at 1200–1599, 13 at 1600–1999, 22 at 2000–2399, and 35 rated 2400+—so RQ3’s difficulty analysis spans the full competitive range rather than only easy tasks. Each C++ submission is compiled with GNU g++ (gnu++17) and run under an OS sandbox (no network, bounded CPU/memory/stack/file size, process-group timeouts) on public and private tests; output comparison is whitespace-tokenized, case-insensitive, 10^{-4} float tolerance. Each run yields one of five verdicts (Table II); when a submission fails differently on different tests we report the most severe under the precedence $CE > RE > TLE > WA > AC$ (a program that does not compile cannot be timed, one that crashes cannot be judged for output, and so on), cached by (problem,

TABLE II
SANDBOX VERDICTS. A *bug* is ANY NON-AC SUBMISSION; THE VERDICT IS THE MOST SEVERE OUTCOME OVER ALL TESTS.

Verdict	Meaning
AC	Accepted: correct output within limits on every test
WA	Wrong Answer: output mismatches the expected output
TLE	Time-Limit Exceeded: exceeds the problem’s CPU budget
RE	Runtime Error: crash, non-zero exit, or sandbox fault
CE	Compile Error: source fails to build with g++

TABLE III
BUG CORPUS OVER 107 PROBLEMS. THE SUCCESS METRIC IS NOT COMPARABLE ACROSS HUMAN AND AI: HUMAN IS A CURATED CORRECT:INCORRECT RATIO; AI IS PER-SAMPLE PASS RATE.

Source	Confirmed bugs	Success metric
Human	8,648	53.4%
gpt-3.5 clean	2,076	3.0%
gpt-5.4-nano modern	1,465	31.5%
gpt-3.5 reflect	1,993	—
Total	14,182	

source, tests). A bug is any submission whose verdict is not AC.

Human solutions come labeled in CodeContests; we re-judge all and keep only label-consistent ones, yielding **8,648 human bugs** (and 9,925 confirmed-correct). **AI** solutions are generated zero-shot from a PaR-style prompt [2] (task framing plus the verbatim statement, which carries the I/O format and examples) at temperature 1.0 with a *fixed* budget of 20 samples per problem (no early stop): gpt-3.5-turbo-0125 (clean) gives 64 AC and **2,076 bugs**; gpt-5.4-nano (modern) gives 675 AC and **1,465 bugs**. We stress that gpt-5.4-nano is *not* a cutoff-controlled source: its training window covers these problems, so it cannot isolate reasoning from memorization. It is included for a different purpose—to additionally check how a *modern, high-capability* model behaves on the same problems—and gpt-3.5-turbo-0125 alone carries the cutoff-aligned claims. Every cutoff-aligned taxonomy result in this paper rests on this proxy model; the modern model is a present-day performance reference, read always against that caveat. **Self-reflection** applies up to three Self-Refine rounds [38] per clean-model bug, continuing the same conversation with the public example-test results, stopping at AC; **1,993 bugs** survive. The total is **14,182 sandbox-confirmed bugs** (Table III).

a) *Scope of the comparison:* The human and AI corpora are not interchangeable behavioral samples. Human submissions are curated CodeContests artifacts with separate correct and incorrect pools; AI submissions are fixed-budget samples from a prompt. We therefore do not use the human correct:incorrect ratio and AI per-sample pass rate to compare programmer ability. The taxonomy comparison is narrower and conditional: *among sandbox-confirmed non-AC C++ submissions on the same problems*, how are bug mechanisms distributed? Acceptance rates are used only for




Fig. 1. Overview of the study pipeline: CodeContests problems are filtered into a post-cutoff C++ benchmark; human, zero-shot AI, and self-reflection submissions are judged by the same sandbox; taxonomy labels come from a fixed Wei et al. taxonomy through a blind human gold set and an audited LLM judge.

within-source trends (e.g., by difficulty) and for the clean-vs-modern contamination/capability contrast. This estimand does not claim that a curated human incorrect submission and a zero-shot LLM sample arise from the same behavioral process.

B. Prompts and protocols

We summarize the three prompt-driven protocols; full prompts are released. *Generation* follows the PaR template [2]: a system message frames the task (“write a correct, efficient C++ program; return only a code block”), and the user message is the verbatim problem statement, which on Codeforces already contains the I/O specification and worked examples; no unit tests or scaffolding are added, so the model must reason from the statement alone. The fixed-budget choice—exactly 20 samples per problem, no early stop—equalizes effort so that a model’s bug yield reflects its true failure rate rather than a collection cap; the clean model yields ≈ 19 bugs per problem, the modern model ≈ 14 . *Self-reflection* continues the same conversation: we append the model’s own buggy program as the assistant turn, then a user turn reporting its results on the public example tests (the input, the expected output, and the program’s actual output for each), and ask for a corrected program, for up to three rounds, stopping at the first AC. The model thus sees exactly the feedback a contestant reads from the sample tests, and nothing more. *Judging* gives gpt-5.4 the problem statement, the buggy code, and the verdict, and asks for the applicable taxonomy leaves plus a one-line rationale in strict JSON, sampled three times with majority aggregation. All model calls use the OpenAI Chat Completions API through the model strings named above; generation/reflection use temperature 1.0, the judge uses temperature 0.4, and no top_p, frequency, or

presence penalties are set. Responses are cached by model, temperature, prompt, and nonce; because proprietary model strings are not immutable snapshots, the released raw generations, judge outputs, and cache keys define the audit record. The API calls that produced the released generation/reflection corpus were completed on 2026-06-29, and the judge labels on 2026-06-30. The cache stores raw response text rather than provider response identifiers, so exact re-generation is not promised; reproducibility means re-running the analysis from the released outputs. The judge is provenance-blind and receives neither the failing hidden test nor a reference solution, matching the human annotators’ information exactly so that disagreement reflects judgment, not evidence. This is a deliberate choice for an apples-to-apples agreement study, but we acknowledge its cost: inferring *which* algorithmic family is at fault for a wrong answer from code alone is genuinely hard, and that irreducible diagnostic noise is itself a plausible contributor to the fine-grained instability we document in RQ1. We therefore keep the common-evidence protocol for the full corpus, and add a separate first-failing-test audit to measure how much stronger evidence changes labels.

C. Labeling and analysis

Labeling is multi-label into the 32 leaves [2]. We size the gold by Cochran’s formula with finite-population correction [45], [46]: 95% confidence, $\pm 5\%$ margin on 14,182 items gives $n \approx 375$. We allocate it arm-balanced (125 each: human, clean, modern), stratify by verdict so rare verdicts appear, and cap four items per problem. **Two annotators** labeled all 375 items provenance-blind (problem statement and buggy code only); a third adjudicated; we report their inter-annotator agreement. The **LLM judge** (gpt-5.4) labels

all 14,182 bugs from the same evidence (statement, code, verdict) with self-consistency over three samples. For verdict-based questions (acceptance, difficulty, repair) we use the exact full corpus; for taxonomy-distribution questions we use the human gold and validate the judge against it (RQ1). We report a descriptive χ^2 /Cramér’s V on family *occurrences*: each multi-label bug contributes one count to every family present. Nine families appear at least once in the zero-shot corpus (GE1–GE6, AE1, AE2, AE5), giving $df = 8$ for pairwise source comparisons; AE3/AE4/AE6 are absent and the GE3 singleton is kept in the test but omitted from the compact table. We do not use this anti-conservative test for inference. Instead, we use a problem-level bootstrap because bugs are multi-label and clustered within problems. For each replicate we sample 107 problem IDs with replacement; within each sampled problem and source, we compute the share of that source’s non-AC bugs whose labels contain each family, skipping problem-arm cells with no bugs; then we average the AI–human difference over sampled problems. We report percentile 95% intervals from 2,000 replicates. This estimates an equal-problem conditional bug-profile difference, rather than letting high-volume problems dominate the comparison.

Three design choices in this protocol are worth making explicit, because they determine what the comparison can support. First, the gold is *arm-balanced* rather than proportional: drawing 125 bugs from each source—not 375 in proportion to the 8,648/2,076/1,465 corpus split—trades population-level precision for equal power to characterize each source’s profile, which is what a human-vs-AI contrast needs. Second, *verdict stratification with rare-verdict floors* guarantees that compile, runtime, and time-limit failures appear in the gold despite being a few percent of the corpus, so the rare-but-diagnostic families (syntax GE5, and the structural failures central to RQ4) are estimable at all. Third, labeling is *provenance-blind by construction*: the annotator and judge packages contain the statement, the code, and the sandbox verdict, with the author identity stripped and the within-problem order shuffled, so neither a human nor the judge can infer “this looks like AI code” and label to a stereotype—the comparison measures the bug, not a guess about its origin. Statistically, we treat the judge as an instrument whose error we bound against the human ceiling (RQ1) rather than as ground truth, and we report effect sizes alongside p -values throughout because, at $n=125$ per arm and with within-problem clustering, significance and magnitude carry different and complementary information.

Finally, after the main analysis, two annotators independently re-labeled a stratified 120-item wrong-answer audit set (40 per zero-shot source) with stronger evidence: the same statement and code plus the first failing test’s input, expected/actual output, and one accepted reference implementation. If a program failed a public sample first, that public sample is the first failing test; otherwise the evidence is private-test based (95 public, 25 private overall). The annotators adjudicated all disagreements, with a third tie-breaker for three items. This audit does not relabel the full corpus; it estimates how far common-evidence diagnostic labels are from root-cause labels

TABLE IV
LABELING AGREEMENT ON THE 375-BUG COMMON-EVIDENCE GOLD (COHEN’S κ , MICRO- F_1).

Pair	κ (leaf)	κ (family)	micro- F_1
human1 vs. human2 (ceiling)	0.80	0.83	0.80
judge vs. human1	0.75	0.77	0.76
judge vs. human2	0.77	0.81	0.77

TABLE V
FIRST-FAILING-TEST/REFERENCE-SOLUTION ROOT-CAUSE AUDIT ON 120 ZERO-SHOT WA SUBMISSIONS (40 PER SOURCE). “CHANGED” COMPARES THE ADJUDICATED ROOT-CAUSE LABEL WITH THE ORIGINAL COMMON-EVIDENCE LABEL.

Metric	Count	Value
Leaf exact agreement	84/120	70.0%
Family exact agreement	85/120	70.8%
Primary-family κ	—	0.69
Binary family κ	—	0.71
Adjudicated disagreements	36/120	30.0%
Leaf changed	70/120	58.3%
Family changed	66/120	55.0%

under stronger evidence.

IV. RESULTS

A. RQ1: How stable are taxonomy labels?

The first result is about the instrument, not the model. The two human annotators on the 375-bug common-evidence gold agree strongly—Cohen’s $\kappa = 0.80$ at leaf level and 0.83 at family level—matching the high agreement reported for this taxonomy by Wei et al. [2]. The `gpt-5.4` judge reaches substantial agreement with both annotators (family $\kappa = 0.77/0.81$, micro- $F_1 \approx 0.77$; Table IV), so it is usable for coarse full-corpus screening. It is not neutral enough for fine distributional claims. On the same 375 gold items the judge assigns the broad design family to 26/34/41% of human/clean/modern bugs, while the human annotators assign 20/28/32%: a 6–9 point inflation that is arm-dependent and largest on the modern arm, which shares the judge’s model family.

The stronger test is evidence stability. In the 120-item audit, annotators relabeled wrong-answer submissions with the first failing test, actual/expected output, and an accepted reference implementation. They again reached substantial agreement (primary-family $\kappa = 0.69$; binary multi-label family $\kappa = 0.71$), then adjudicated all 36 disagreements. Against the original common-evidence labels, the adjudicated root-cause labels changed 70/120 leaves (58.3%) and 66/120 families (55.0%; Table V). The rate is not confined to one source: family labels changed for 65.0% of human submissions and 50.0% of both AI arms.

This is the construct boundary for the rest of the paper. Common-evidence labels can say that a failure looks mathematical, design-like, or syntactic under the evidence a scalable study can usually afford. They cannot, by themselves, establish

the root-cause mechanism of a wrong answer. Therefore, when a full-corpus family shift is only a few percentage points, judge granularity and missing failing-test evidence can dominate the effect. We treat such shifts as exploratory screens and reserve strong claims for coarse patterns and label-free verdict dynamics.

B. RQ2: What bug-profile claims remain defensible?

On the human gold (Table VII) all three sources are dominated by the same two families—mathematics (AE1, $\approx 33\%$) and design (GE1, 20–32%)—and their high-level general-vs-algorithmic split is similar without being diagnostic. A *weak* directional tendency is visible: design (GE1) runs 20% (human) / 28% (clean) / 32% (modern), while implementation boundaries (GE2: 11/7/5%) and I/O formatting (GE6: 6% vs. $< 1\%$) run the other way. But the tendency is, on the gold, *underpowered*. At 125 bugs per arm the three-way family χ^2 can detect only $V \gtrsim 0.26$ at 80% power (a minimum-detectable-effect we compute by inverting the noncentral χ^2), and the observed $V \approx 0.23$ ($p \approx 0.07$ – 0.10 per annotator) sits just below that floor—so the gold can neither confirm the tendency nor rule it out, and we do not read its non-significance as evidence of no effect.

a) *Full-corpus shifts are small, not self-interpreting*: The full corpus helps with precision but not automatically with validity. A descriptive χ^2 on the source $\times$ family occurrence table over the nine observed families (Sec. III; $df = 8$) gives high significance at a negligible effect size—Cramér’s $V = 0.11$ (human–clean, $p = 3 \times 10^{-25}$), 0.08 (human–modern)—and the problem-level bootstrap does detect several small shifts (Table VI). The clean model has fewer design labels than humans (-5.6 points) and more mathematics/syntax labels ($+4.0/+2.6$); the modern model shows smaller shifts, with syntax and mathematics up by about 2–3 points and boundary/I/O down by about 1–2 points. These are not large taxonomy-defining separations: all major shifts are within the 6–9 point design inflation that RQ1 finds in the judge itself, and most are much smaller. We therefore use the judge-labeled full corpus as a *screen* for fine shifts, not as a causal claim that humans and AI have different diagnostic-label distributions. The robust RQ2 result is the coarse one: all sources concentrate in mathematics and design, and low-level coding families remain minority modes.

b) *The apparent reversal is a sample effect, not a labeler effect*: A natural worry is that the comparison reverses under the judge: design runs 20/28/32 (human/clean/modern) under the human annotators but 51/45/52 under the judge on the full corpus, the clean model dropping to last. Holding the *sample* fixed dispels this (Fig. 2). On the same 375 gold items the judge gives 26/34/41—it *inflates* design by ~ 6 –9 points (a real granularity bias, quantified in RQ1) but *preserves* the human $<$ clean $<$ modern ordering. The reversal appears only when the judge moves from the stratified gold to the natural corpus (26/34/41 $\rightarrow$ 51/45/52), so it is driven by the *sample* (verdict-stratified, per-problem-capped gold vs. the full corpus), not the labeler. Neither sample nor labeler yields




Fig. 2. Deconfounding the apparent design (GE1) reversal. (a) Holding the *sample* fixed (the same 375 gold items), the LLM judge inflates the design family ($\sim +6$ –9 points) but *preserves* the human $<$ clean $<$ modern ordering—its granularity bias is a level shift, not a reversal. (b) Holding the *labeler* fixed (judge), moving from the stratified gold to the full corpus changes the ordering (clean drops to last): the “reversal” is a sample effect. Neither route yields a robustly ordered difference.

a stable ordered design story, which is why we treat the fine shifts as measurement-limited rather than explanatory.

Concrete cases match the quantitative result. On problem 1619B, a human and the clean model both use the right inclusion–exclusion idea for counting squares/cubes but both lose to floating-point root precision (AE1.1). Divergent cases exist—on 1608A the AI misreads the required array while a human solution only mishandles the test-case harness—but they are examples of the fine texture that RQ1 makes hard to measure reliably, not a stable distributional separation.

c) *Leaf-level texture is hypothesis-generating*: The gold still helps explain what the coarse families mean. The single most common leaf everywhere is a mathematical formula error (AE1.1, $\approx 33\%$), shared rather than discriminative. Some runner-up leaves lean differently—AI bugs more often look like requirement or strategy mistakes, while human bugs include harness and DP transition slips—but this is exactly the leaf-vs-family granularity RQ1 flags as unstable. We therefore treat leaf examples as qualitative texture, not as a measured distributional claim.

The one structural regularity worth stating is what is *shared*: across both the gold and the full corpus, and across all three sources, the coding-level families that a linter or compiler would catch—syntax (GE5) and data type (GE4)—sit at similar, low rates, while the mass concentrates in the two “thinking” families (mathematics AE1, design GE1). In other words, where the sources are alike is not incidental: modern code generation has largely closed the low-level coding gap, leaving everyone—human and model—to fail mostly on the algorithmic substance. That convergence, rather than any residual tilt, is the robust RQ2 result.

This coarse result is not just an artifact of verdict mix or near-impossible clean-model attempts. In the WA body alone (compile failures removed), math+design covers 73.8% of human bugs, 70.4% of clean-model bugs, and 73.6% of modern-model bugs; low-level data/syntax/I/O families fall below 2.4% in every arm. Averaging AI–human differences over equal problem $\times$ verdict cells gives a clean-model math+design difference of only -0.6 points (cell percentile in-

TABLE VI

FULL JUDGE-LABELED CORPUS. PERCENT COLUMNS ARE CORPUS SHARES OF BUGS WITH A FAMILY PRESENT. Δ COLUMNS ARE EQUAL-PROBLEM AI-HUMAN DIFFERENCES WITH PROBLEM-BOOTSTRAP 95% CIs (2,000 REPLICATES); PROBLEM-ARM CELLS WITH NO BUGS ARE UNDEFINED FOR THE CONDITIONAL BUG-PROFILE ESTIMATE. THE SHIFTS ARE STATISTICALLY VISIBLE BUT SMALL RELATIVE TO THE JUDGE GRANULARITY BIAS MEASURED IN RQ1.

Family	Human	Clean	Modern	Δ Clean-H	Δ Modern-H
GE1 Design	51.4	44.8	52.3	-5.6[-7.9, -3.4]	-0.7[-3.9, +2.3]
GE2 Boundary	8.2	9.9	6.3	-0.6[-2.0, +1.0]	-2.3[-4.0, -0.4]
GE4 Data type	1.0	0.1	0.1	-0.7[-1.5, -0.1]	-0.7[-1.5, -0.1]
GE5 Syntax	2.6	6.1	5.1	+2.6[+0.8, +4.6]	+2.9[+0.7, +5.4]
GE6 I/O	1.2	0.0	0.1	-1.5[-2.0, -1.0]	-1.3[-1.9, -0.7]
AE1 Mathematics	19.4	19.9	17.4	+4.0[+2.3, +5.8]	+2.4[+0.9, +4.0]
AE2 Greedy	7.7	9.2	10.1	+0.4[-0.4, +1.3]	+0.8[-0.0, +1.5]
AE5 DP	8.5	10.1	8.7	+1.3[+0.4, +2.4]	-0.9[-2.9, +0.7]
Cramér's V : H-C 0.11, H-M 0.08, C-M 0.09					

TABLE VII

BUG-FAMILY DISTRIBUTION ON THE HUMAN GOLD (% OF LABELS, MEAN OF TWO ANNOTATORS); BOLD MARKS THE LARGER HUMAN/AI SIDE. THIS IS A HUMAN-LABELED GOLD-SAMPLE TENDENCY, NOT A DEFINITIVE POPULATION ESTIMATE.

	Family	Human	Clean	Modern
GE1	Design / wrong algorithm	20	28	32
GE2	Boundary / edge cases	11	7	5
GE5	Syntax	9	8	9
GE6	Input/Output	6	1	1
AE1	Mathematics	32	35	33
AE2	Greedy	8	13	9
AE5	Dynamic programming	8	6	8

terval $[-17.5, +10.5]$), and WA-only equal-problem averaging gives -0.9 points. Finally, restricting to the 14 problems where the clean model produces at least one accepted solution makes both humans and the clean model more mathematics-heavy, but the profile remains shared: math+design covers 82.0% of human bugs and 80.6% of clean bugs. These checks reinforce the scope limit, but they do not overturn the coarse conclusion.

C. RQ3: Capability, contamination, and difficulty

Table III reports three success metrics, but they answer different questions: the human number is a curated correct:incorrect ratio, whereas the AI numbers are fixed-budget per-sample pass rates. We therefore do not use the table as a cross-source ability comparison. The defensible signal is the *within-source* trend with difficulty (Table VIII, Fig. 3). The clean model is at the floor everywhere ($\leq 8\%$). The modern model is much stronger but not simply solved by lookup: it falls from 55.0% on the easiest band to 12.3% on the hardest, with a non-monotone middle band (21.2% then 35.0%) that points to problem-set heterogeneity rather than a smooth ability curve. Its high acceptance on some hard problems is consistent with partial contamination, but the overall decline argues against pure memorization. The human rate is flat (42–59%), a survivorship artifact: harder problems attract stronger solvers.

The contrast between the two models is the point of the contamination design. The clean model, whose cutoff predates every problem, never rises above 8% acceptance and is flat

TABLE VIII

PER-SUBMISSION ACCEPTANCE (%) BY DIFFICULTY BAND.

Difficulty	Human	Clean	Modern
≤ 1199	59.3	7.6	55.0
1200–1599	49.5	0.8	43.8
1600–1999	42.9	0.4	21.2
2000–2399	55.5	3.6	35.0
2400+	53.3	1.0	12.3

across difficulty—the behavior expected of a 2023-class model attempting unseen contest tasks from the statement alone. The modern model, whose training window covers the problems, scores far higher and declines overall, though not monotonically, with difficulty; this profile is hard to reconcile with pure lookup. Two readings are consistent with the data—real capability gains since the clean model, and partial memorization of the easier, more-discussed problems—and we do not claim to separate them, only to show that an aggregate pass rate would conflate them whereas a difficulty-resolved one begins to pull them apart. The clean model remains a cutoff-aligned proxy anchor for distributional claims; the modern model is the contrast that makes the anchor's value legible.

The verdict mix is exact and label-free (Table IX). All sources are wrong-answer dominated (90–95%), but AI fails to *compile* more often than humans (clean 6.1%, modern 5.1% vs. 2.6%) and the clean model crashes more (RE 2.6%); the modern model has the fewest runtime errors (0.5%), consistent with cleaner low-level coding as capability grows. That compile failures survive at all is itself notable: a 6% compile rate under a fixed prompt means the weak model periodically emits C++ that does not build, a failure mode essentially absent from curated human submissions and one that the self-reflection loop only partly clears (RQ4).

D. RQ4: What does self-reflection repair?

We study one concrete, common setup—the clean model revising its own code with *public example-test* feedback for three rounds—and report what it can and cannot fix; the numbers below are specific to this model and this feedback signal, not a claim about self-reflection in general. Part of why the rate is so low is that the public examples were already in

TABLE IX
VERDICT DISTRIBUTION PER SOURCE (% OF CONFIRMED BUGS, EXACT).

Source	WA	TLE	RE	CE
Human	94.6	0.9	1.8	2.6
Clean	90.0	1.3	2.6	6.1
Modern	93.5	0.9	0.5	5.1




Fig. 3. Acceptance vs. difficulty: the clean model is at the floor, the modern model declines overall but not monotonically, and humans are flat (survivorship).

the original prompt (they ship with the Codeforces statement); reflection’s *new* information is therefore only the model’s own wrong output on inputs it had already seen, not fresh test data—so it can correct a mis-handled visible case but is blind to everything the hidden tests check. In that setting, repair reaches only **4.0%** (83/2,076) of the clean model’s bugs, and that number is not distinguishable from a trivial same-budget alternative. Using the existing 20-sample zero-shot budget as a leave-one-out pool, three fresh draws from the other 19 cached samples for each failed attempt would fix **82.6** bugs in expectation (**4.0%**); the difference from reflection is below 0.1 points. Table X therefore changes the interpretation of RQ4: public-example reflection is not an efficient repair strategy over simple resampling; it is a diagnostic probe for which bugs survive weak feedback. A post-hoc edit audit of the 83 fixed trajectories agrees: 7 fixes are compile-error repairs, 16 leave the original program at least 95% character-identical, 47 occur on ≤ 900 -rated problems, and three problems account for 46 fixes. Figure 4 tracks every bug’s verdict across the three rounds: the accepted band grows $0 \rightarrow 53 \rightarrow 73 \rightarrow 83$, so most repairs land in the first round and the third adds only ten, while the wrong-answer mass ($\approx 1,850$) is essentially static. The rate is sharply stratified by verdict: compile errors reach 5.5% and wrong answers 4.1%, but **runtime (0/55) and time-limit (0/27) errors are never repaired**—runtime errors even tick up as the model perturbs code while chasing a wrong answer. Structural failures—a wrong-complexity algorithm, a crash on an unseen input—pass the small public examples unchanged, so reflection sees no signal. Because RQ1 validates automated labels only at family level, Fig. 5 uses the judge’s primary family labels over all 2,076 clean-model bugs. The result is

TABLE X
SELF-REFLECTION VS. SAME-BUDGET RESAMPLING FOR CLEAN-MODEL BUGS. RESAMPLING ROWS ARE LEAVE-ONE-OUT EXPECTATIONS FROM EACH PROBLEM’S EXISTING 20 CACHED ZERO-SHOT SAMPLES, NOT NEW MODEL CALLS.

Strategy	Extra calls	Fixed	Rate
Self-reflection	3 repair	83	4.0%
Fresh sample	1	37.6 exp.	1.8%
Fresh samples	2	63.3 exp.	3.0%
Fresh samples	3	82.6 exp.	4.0%




Fig. 4. Verdict transitions of all 2,076 clean-model bugs across the three self-reflection rounds. The accepted (green) band grows only $0 \rightarrow 53 \rightarrow 73 \rightarrow 83$; wrong answers dominate and persist, and runtime/time-limit failures never resolve.

still narrow: repairs concentrate in easy syntax/design cases, while boundary, mathematical, greedy, and DP failures mostly persist. The surviving diagnostic-label profile is therefore almost unchanged (general-error share $61\% \rightarrow 61\%$): the AI’s failure signature is *robust to self-correction*.

a) *Repair power collapses with difficulty*: Splitting the same trajectories by problem difficulty and original bug family (Fig. 5) shows the repair rate falling monotonically with rating: **7.1%** on easy problems (≤ 1599), **3.4%** on medium (1600–2399), and **1.4%** on hard (2400+). On hard problems the accepted band barely moves after the first round ($8 \rightarrow 10 \rightarrow 10$ of 693), whereas easy problems keep accruing fixes ($30 \rightarrow 43 \rightarrow 50$ of 700). Reflection thus helps exactly where the model was already close—easy problems with surface slips—and is nearly inert on the hard, design-bound problems where its bugs concentrate (RQ2, RQ3). Self-correction is not a substitute for choosing the right algorithm. Two secondary dynamics point the same way: compile errors—the one locally diagnosable failure—are cleared in every band (e.g. $54 \rightarrow 26$ on hard), but on hard problems they convert to *wrong answers* rather than acceptances; runtime errors slightly *increase* on medium and hard problems as edits chasing a wrong answer introduce new crashes.

V. DISCUSSION

a) *Alike where it counts*: The measurement result is the main point: on identical problems, the human and AI bugs we measure occupy the same dominant families, but the remaining

Original bug family	Repair rate by difficulty			All 83/2076 fixed
	Easy	Medium	Hard	
GE1 Design	11.4% 31/272	5.5% 16/291	2.7% 10/366	6.1% 57/929
GE2 Boundary	1.1% 1/92	0.0% 0/70	0.0% 0/43	0.5% 1/205
GE5 Syntax	20.6% 7/34	0.0% 0/39	0.0% 0/54	5.5% 7/127
AE1 Math	3.4% 8/237	3.7% 3/81	0.0% 0/95	2.7% 11/413
AE2 Greedy	4.7% 3/64	3.4% 3/89	0.0% 0/37	3.2% 6/190
AE5 DP	- 0/0	0.9% 1/113	0.0% 0/97	0.5% 1/210
Low	High repair rate			

Fig. 5. Taxonomy-aware repair rate of self-reflection, using the original bug’s AI-judge primary family label after family-level validation against the 375-item human gold. Cells show fixed/initial counts and percentages by difficulty; rows show families with at least 100 initial clean-model bugs (two singleton families omitted). Repairs concentrate in easy syntax/design cases and are nearly absent for hard algorithmic families.

family-level shifts are small enough that judge granularity, sample design, and clustering matter as much as the effects themselves.

b) Public-example retry is not a safety net: In our weak-feedback setting, reflection never repairs structural failures and barely shifts the diagnostic-label profile. Pipelines that rely on retrying after public examples need stronger external signals—adversarial tests, complexity analysis, or human review—for the strategic failures that dominate this corpus.

c) Evaluation and measurement: Pass@ k is not enough once the remaining failures are algorithmic and small tests miss them. A taxonomy-grounded profile is a useful complementary axis, but RQ1 shows that it must be read carefully: a substantial-agreement judge can still be too coarse to adjudicate small distributional differences, and first-failing evidence can change the family label entirely. Reliability must be established at the granularity of the claim, not only by an aggregate κ .

VI. THREATS TO VALIDITY

We organize threats following standard empirical-SE guidance [52], [53].

Construct validity. The human acceptance “rate” is the correct:incorrect ratio of CodeContests’ curated submissions, not a real submission success rate; we therefore compare only within-source *trends* with difficulty rather than absolute rates across sources. Taxonomy labeling is inherently subjective; we mitigate with two independent provenance-blind annotators and third-party adjudication, and we report inter-annotator $\kappa = 0.80$ as the reliability ceiling. The full-corpus annotators and judge do not receive the failing test or a reference solution, so those labels are diagnostic hypotheses under common evidence, not demonstrated root causes. Our 120-item root-cause audit quantifies the gap: first-failing-test/reference evidence changes 55.0% of family labels. As RQ1 shows, the construct is sensitive to both evidence and

labeling granularity—an over-broad “wrong algorithm” default collapses the contrast—so we fixed an explicit “prefer the most specific leaf” convention and avoid fine causal claims. A genuine, unresolved threat is that the judge (gpt-5.4) and the modern model (gpt-5.4-nano) share a model family. We do not defend this with an aggregate κ —RQ1’s whole point is that an aggregate κ cannot see a directional granularity bias—and indeed, at the granularity level, the judge’s excess design labels (those it adds where humans do not) grow across arms, 6%/8%/10% for human/clean/modern, peaking on the same-family modern arm (Sec. IV-A). We cannot separate “the modern arm genuinely has more design bugs” from “the judge over-recognizes its own family’s output,” and we flag it as a limitation rather than explain it away; it is one more reason we rest no fine distributional claim on judge labels.

Internal validity. Excluding generated tests for tractability admits occasional weak-test pass-through, which can mislabel a buggy program as correct; this affects all sources symmetrically and so should not bias the human–AI comparison. Verdicts and acceptance come from a deterministic, cached sandbox and are not subject to labeling error. The modern model’s contamination is uncontrolled by design, so its results are a contrast, not a clean measurement. Contamination can also change the *conditional* diagnostic-label profile: if memorized or easy-to-recall solutions become AC, they leave the non-AC pool whose labels we analyze. This selection effect is another reason we make no taxonomy claim from the modern arm; the clean model anchors only a cutoff-aligned proxy signal.

External validity. We study C++ Codeforces problems, one PaR-style prompt at temperature 1.0, and two models (one cutoff-aligned, one modern); this single language, prompt, and clean proxy model bound the scope. The clean-model limit is *structural*: post-2021 problems require a pre-2021 training cutoff, which rules out most stronger code models

and is the same tension that motivates contamination-aware benchmarks [7]. Even within this limit, our scope exceeds the closest prior taxonomy study, which labeled 75 programs from one model with no human baseline [2]. The clean model’s 3% acceptance is therefore a property of this realistic hidden-test setting, not a general statement about current LLMs. We claim only coarse similarity plus measurement-limited fine shifts; a different language, prompt, or stronger clean model could surface larger differences.

Conclusion validity. Bugs cluster within problems and labels are multi-label, so a χ^2 that treats the 14,182 bugs as independent is anti-conservative; we therefore base inference on a bootstrap over the 107 problems (the true unit), reported explicitly in Table VI. Some bootstrap intervals exclude zero, but the effect sizes are small and estimated with the same judge whose family-level bias is 6–9 points on the gold. We are equally explicit about power: at 125 bugs/arm the gold detects only $V \gtrsim 0.26$, so its borderline tendency ($V \approx 0.23$) is *underpowered*—we do not read its non-significance as proof of no effect. These legs point to a bounded conclusion: coarse profiles are stable, while fine human–AI differences remain measurement-limited rather than confirmed causal separations.

VII. CONCLUSION

On identical post-cutoff problems under our cutoff-aligned proxy design, the human and AI failed submissions we measure share the same coarse diagnostic-label profile: mathematics and design dominate, while low-level coding failures remain minority modes. The fine profile is not settled. Full-corpus judge labels reveal small family shifts, but the human gold is underpowered, the judge’s granularity bias is of the same order as those shifts, and our first-failing-test audit shows that stronger evidence changes 55% of family labels. That is the paper’s central methodological lesson: comparing LLM and human diagnostic-label *distributions* demands evidence, labeler validation, sample control, and clustering control at the granularity of the claim, not merely an aggregate κ . What is unambiguous is *within-source dynamics*: the clean proxy stays near the floor, the modern contaminated reference is stronger but difficulty-sensitive, and self-reflection with public-example feedback repairs only 4% of the clean model’s bugs—no better than a same-budget resampling baseline—and never a runtime or time-limit failure. Human and AI bugs look similar at the coarse resolution current measurement can defend; pinning down the rest will require sharper instruments.

VIII. ARTIFACT AVAILABILITY

All artifacts are available at <https://anonymous.4open.science/r/aitaxo>.

REFERENCES

- [1] Y. Li *et al.*, “Competition-level code generation with AlphaCode,” *Science*, vol. 378, no. 6624, pp. 1092–1097, 2022.
- [2] M. Wei, Z. Li, X. Chen, M. Zheng, Z. Qu, C. Yu, S. Chen, and X. Ju, “Evaluating and improving LLM-based competitive program generation,” *Information and Software Technology*, vol. 191, p. 107977, 2026.
- [3] M. Chen *et al.*, “Evaluating large language models trained on code,” *arXiv preprint arXiv:2107.03374*, 2021.
- [4] J. Austin *et al.*, “Program synthesis with large language models,” *arXiv preprint arXiv:2108.07732*, 2021.
- [5] D. Hendrycks *et al.*, “Measuring coding challenge competence with APPS,” in *NeurIPS Datasets and Benchmarks Track*, 2021.
- [6] J. Liu, C. S. Xia, Y. Wang, and L. Zhang, “Is your code generated by ChatGPT really correct? Rigorous evaluation of large language models for code generation,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [7] N. Jain *et al.*, “LiveCodeBench: Holistic and contamination-free evaluation of large language models for code,” *arXiv preprint arXiv:2403.07974*, 2024.
- [8] OpenAI, “GPT-4 technical report,” *arXiv preprint arXiv:2303.08774*, 2023.
- [9] A. Vaswani *et al.*, “Attention is all you need,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [10] T. B. Brown *et al.*, “Language models are few-shot learners,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [11] E. Nijkamp *et al.*, “CodeGen: An open large language model for code with multi-turn program synthesis,” in *International Conference on Learning Representations (ICLR)*, 2023.
- [12] B. Rozière *et al.*, “Code Llama: Open foundation models for code,” *arXiv preprint arXiv:2308.12950*, 2023.
- [13] H. Touvron *et al.*, “LLaMA: Open and efficient foundation language models,” *arXiv preprint arXiv:2302.13971*, 2023.
- [14] D. Fried *et al.*, “InCoder: A generative model for code infilling and synthesis,” in *International Conference on Learning Representations (ICLR)*, 2023.
- [15] Y. Wang, W. Wang, S. Joty, and S. C. Hoi, “CodeT5: Identifier-aware unified pre-trained encoder-decoder models for code understanding and generation,” in *Proceedings of EMNLP*, 2021.
- [16] D. Guo *et al.*, “DeepSeek-Coder: When the large language model meets programming,” *arXiv preprint arXiv:2401.14196*, 2024.
- [17] X. Hou *et al.*, “Large language models for software engineering: A systematic literature review,” *ACM Transactions on Software Engineering and Methodology*, 2024.
- [18] T. Y. Zhuo *et al.*, “BigCodeBench: Benchmarking code generation with diverse function calls and complex instructions,” *arXiv preprint arXiv:2406.15877*, 2024.
- [19] C. E. Jimenez *et al.*, “SWE-bench: Can language models resolve real-world GitHub issues?” in *International Conference on Learning Representations (ICLR)*, 2024.
- [20] S. Lu *et al.*, “CodeXGLUE: A machine learning benchmark dataset for code understanding and generation,” in *NeurIPS Datasets and Benchmarks Track*, 2021.
- [21] R. Puri *et al.*, “Project CodeNet: A large-scale AI for code dataset for learning a diversity of coding tasks,” in *NeurIPS Datasets and Benchmarks Track*, 2021.
- [22] S. Kulal *et al.*, “SPoC: Search-based pseudocode to code,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [23] O. Sainz *et al.*, “NLP evaluation in trouble: On the need to measure LLM data contamination for each benchmark,” in *Findings of EMNLP*, 2023.
- [24] S. Golchin and M. Surdeanu, “Time travel in LLMs: Tracing data contamination in large language models,” in *International Conference on Learning Representations (ICLR)*, 2024.
- [25] M. Riddell, A. Ni, and A. Cohan, “Quantifying contamination in evaluating code generation capabilities of language models,” *arXiv preprint arXiv:2403.04811*, 2024.
- [26] I. Magar and R. Schwartz, “Data contamination: From memorization to exploitation,” in *Proceedings of ACL*, 2022.
- [27] R. Chillarege *et al.*, “Orthogonal defect classification—a concept for in-process measurements,” *IEEE Transactions on Software Engineering*, vol. 18, no. 11, pp. 943–956, 1992.
- [28] L. Tan, C. Liu, Z. Li, X. Wang, Y. Zhou, and C. Zhai, “Bug characteristics in open source software,” *Empirical Software Engineering*, vol. 19, no. 6, pp. 1665–1705, 2014.
- [29] R.-M. Karampatsis and C. Sutton, “How often do single-statement bugs occur? The ManySStuBs4J dataset,” in *Proceedings of MSR*, 2020, pp. 573–577.
- [30] G. Catolino, F. Palomba, A. Zaidman, and F. Ferrucci, “Not all bugs are the same: Understanding, characterizing, and classifying bug types,” *Journal of Systems and Software*, vol. 152, pp. 165–181, 2019.

- [31] R. Just, D. Jalali, and M. D. Ernst, "Defects4J: A database of existing faults to enable controlled testing studies for Java programs," in *Proceedings of ISSTA*, 2014, pp. 437–440.
- [32] C. Le Goues *et al.*, "The ManyBugs and IntroClass benchmarks for automated repair of C programs," *IEEE Transactions on Software Engineering*, vol. 41, no. 12, pp. 1236–1256, 2015.
- [33] M. Tufano *et al.*, "An empirical study on learning bug-fixing patches in the wild via neural machine translation," *ACM Transactions on Software Engineering and Methodology*, vol. 28, no. 4, pp. 1–29, 2019.
- [34] H. Pearce, B. Ahmad, B. Tan, B. Dolan-Gavitt, and R. Karri, "Asleep at the keyboard? Assessing the security of GitHub Copilot's code contributions," in *IEEE Symposium on Security and Privacy (S&P)*, 2022, pp. 754–768.
- [35] M. Monperrus, "Automatic software repair: A bibliography," *ACM Computing Surveys*, vol. 51, no. 1, pp. 1–24, 2018.
- [36] L. Gazzola, D. Micucci, and L. Mariani, "Automatic software repair: A survey," *IEEE Transactions on Software Engineering*, vol. 45, no. 1, pp. 34–67, 2019.
- [37] C. S. Xia, Y. Wei, and L. Zhang, "Automated program repair in the era of large pre-trained language models," in *Proceedings of ICSE*, 2023, pp. 1482–1494.
- [38] A. Madaan *et al.*, "Self-Refine: Iterative refinement with self-feedback," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [39] X. Chen, M. Lin, N. Schärli, and D. Zhou, "Teaching large language models to self-debug," *arXiv preprint arXiv:2304.05128*, 2023.
- [40] N. Shinn *et al.*, "Reflexion: Language agents with verbal reinforcement learning," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [41] T. X. Olausson *et al.*, "Is self-repair a silver bullet for code generation?" in *International Conference on Learning Representations (ICLR)*, 2024.
- [42] L. Zheng *et al.*, "Judging LLM-as-a-judge with MT-Bench and Chatbot Arena," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [43] Y. Liu *et al.*, "G-Eval: NLG evaluation using GPT-4 with better human alignment," in *Proceedings of EMNLP*, 2023.
- [44] P. Wang *et al.*, "Large language models are not fair evaluators," in *Proceedings of ACL*, 2024.
- [45] W. G. Cochran, *Sampling Techniques*, 3rd ed. John Wiley & Sons, 1977.
- [46] R. V. Krejcie and D. W. Morgan, "Determining sample size for research activities," *Educational and Psychological Measurement*, vol. 30, no. 3, pp. 607–610, 1970.
- [47] J. Cohen, "A coefficient of agreement for nominal scales," *Educational and Psychological Measurement*, vol. 20, no. 1, pp. 37–46, 1960.
- [48] J. R. Landis and G. G. Koch, "The measurement of observer agreement for categorical data," *Biometrics*, vol. 33, no. 1, pp. 159–174, 1977.
- [49] M. L. McHugh, "Interrater reliability: The kappa statistic," *Biochemia Medica*, vol. 22, no. 3, pp. 276–282, 2012.
- [50] R. Artstein and M. Poesio, "Inter-coder agreement for computational linguistics," *Computational Linguistics*, vol. 34, no. 4, pp. 555–596, 2008.
- [51] K. Krippendorff, *Content Analysis: An Introduction to Its Methodology*, 4th ed. Sage, 2018.
- [52] P. Runeson and M. Höst, "Guidelines for conducting and reporting case study research in software engineering," *Empirical Software Engineering*, vol. 14, no. 2, pp. 131–164, 2009.
- [53] C. Wohlin, P. Runeson, M. Höst, M. C. Ohlsson, B. Regnell, and A. Wesslén, *Experimentation in Software Engineering*. Springer, 2012.